



# NationBrowse

---

Making sense of civic data

## CS4970 Group 2

Graham Greenfield

Jeremy Howard

Nick Roma

Mike Tigas

## Live URL

<http://capstone.nationbrowse.com/>

## Source Repository

[http://github.com/mtigas/cs4970\\_capstone](http://github.com/mtigas/cs4970_capstone)

# Table of Contents

---

<b>List of figures &amp; tables</b>	<b>i</b>
<b>Executive summary</b>	<b>1</b>
<b>Problem definition</b>	<b>1</b>
<b>Requirements analysis</b>	<b>3</b>
System constraints	3
Performance requirements	4
Resource requirements	4
Alternative solutions	4
Evaluation metrics	5
<b>Design Specifications</b>	<b>5</b>
Overview	5
Software design	5
Data requirements	5
Hardware	5
Testing methods	6
Scheduling diagram	6
Implementation resources & costs	6
<b>Product implementation</b>	<b>6</b>
<b>Evaluation of product performance</b>	<b>7</b>
<b>Discussion, future work, and conclusions</b>	<b>7</b>
<b>References</b>	<b>8</b>

## List of figures & tables

---

## Executive summary

---

NationBrowse is a Web application that provides a browsable front-end to civic data provided by the United States Government, through the Data.gov Web site. The primary goal of the Web application is to not only provide an easy-to-use system for browsing statistics, but to also offer statistical analysis and comparison functionality.

The sheer amount of data available from the 2000 United States Census presents a significant problem to those who may be interested in demographic information. (Note that the 2000 Census represents only a fraction of available datasets on Data.gov.) Providing an application interface to assist in analyzing the data would be extremely useful to research organizations, marketers, investigative reporters, and the public at-large.

In addition to statistical analysis, the application will also include map-based browsing of datasets, rendered as a color-coded map of the respective geographic areas — states, counties, ZIP codes, etc.

## Problem definition

---

[Draft note: For this draft, the Introduction, Background, Goals & Objectives, and Approach sections have been combined into one section due to their inherent inter-relatedness. This is subject to change before this draft is finalized.]

The overwhelming amount of information available via Data.gov makes it difficult for people and organizations to use. While releasing large amounts of government data presents a step toward transparency, making the data useful actualizes the goal of the government's Data.gov project.

Census data is not new — the United States Government digitally released the results in 2001 and these datasets have been widely used. However, only recently have Web technologies evolved to the point where a user-centric browsable presentation of the data is possible.

Few dynamic applications to civic data have existed until very recently. Since the launch of the Data.gov project in 2008, Sunlight Labs, a pro-transparency think tank, has sponsored an annual contest for applications (Apps For America) that use civic data from the Data.gov Web site.[1][2] Apps For America recognizes applications of Data.gov datasets that provide a useful tool to the public. Several apps from the Apps For America 2 contest appear to have goals that fall in-line with those of NationBrowse:

- ThisWeKnow[3] provides a browsable interface of many Data.gov datasets, but does not perform any analyses or provide any dynamic element to the browsing experience. There are no maps or alternative perspectives on the data, everything is simply presented as a list of trivia regarding the user's selected location.
- DataMasher[4] provides a way to compare states against one another by “mashing up” two disparate datasets. For example, one may divide crime rate by high school graduation rate, and create a color-coded map that displays that relationship in each state. DataMasher admits that the Web site's output is not statistically rigorous, but is simply a tool to drive conversations regarding the data.

In addition to Web sites featured on Apps For America, several other Web sites have inspired this project either through use of civic data or the use of specific open-source technologies:

- The Los Angeles Times launched a “neighborhood mapping” project[7] that provides lists of demographic information for each L.A. neighborhood, in a very similar vein to the ThisWeKnow project, above.
- EveryBlock[5] is a Web site developed by several former programmers in the news industry, which provides aggregation of local civic data (police records, building permits) in addition to news stories and other information. EveryBlock provides a custom map interface to much of it's data. Paul Smith, a primary developer for EveryBlock's mapping infrastructure, recently wrote about the open-source GIS stack

that EveryBlock uses.[6] It is hoped that a implementation of this “stack” will drive some portions of NationBrowse’s mapping applications.

- 538
- NYT election map
- NYT maps

A quick look at the above sources shows the wide variety of tools available to browse demographic data. However, none of the above products creates any sort of statistical analysis regarding the wide variety of data supplied by the U.S. Government. In addition, no product has fully combined a dynamic GIS mapping application (such as EveryBlock’s) with analyzed civic data. It is our hope that NationBrowse allows a more thorough and useful (in a significant research sense) look into the information that the United States Government has released.

## Requirements analysis

---

### System constraints

As a data-driven Web application, NationBrowse is inherently constrained by the availability and usability of data supplied by the source(s). Because the data will be manually imported into a database internal to the application, updates to the data are also restricted.

The project is intended to be a publicly-accessible web site. Further application constraints come from the server load of the application and it’s respective database.

Some portions of the map renderer may rely on external services, such as Google Maps or other accessible map tile sources.

## Performance requirements

The project must be accessible by modern Web browsers, over HTTP. The server must dynamically serve map tiles to the client — performance overhead regarding this is alleviated by the use of TileCache software.

NationBrowse will require statistical analyses to be performed over various datasets. Some functions (mean, standard deviation) are built-in to the PostgreSQL database. Others are implemented in the SciPy Python library that the project will use, while other algorithms may have to be developed.

## Resource requirements

The application software requires an HTTP Web server with access to Python. The map rendering and GIS portions of the application are built toward the PostgreSQL database using PostGIS extensions.

The data that drives the application comes from the United States Government on Data.gov. Specifically, GIS data will be sourced from the Census Bureau's TIGER/Line shapefiles, which provide detailed shape/border data. Other data (demographics) will be sourced from Census Bureau data and other datasets available on Data.gov.

In lieu of school servers lacking Python and GIS software, the project's development server operates on a Slicehost[8] virtual server, owned by one of the group members.

## Alternative solutions

We chose Django as the overall application framework, because it provides an ORM (object-relational mapping) API that handles much of the database work. From a technical standpoint, the database, Web page, and statistical portions of the application could be implemented in any Web-ready language (such as PHP), but Django and Python provide much in terms of utility, so that many low-level tasks (such as interfacing with the database or serving templated pages) are made simpler. Using

Django and Python allows us to concentrate our efforts on the statistical and mapping portions of the application.

Mapnik was selected as the GIS backend because of it's free and open-source nature and compatibility with other portions of the application — Mapnik utilizes Python, as do other portions of our application. Alternative solutions include MapServer, another free and open-source project, and ArcGIS Server, which is a commercial product.

## Evaluation metrics

As a Web application, the project will mostly be evaluated on feature-completeness and performance metrics. Django and Python offer various methods to test performance: the primary performance metrics would be server response time (or page load time), CPU load, and database load (as measured in number of queries).

# Design Specifications

---

[Draft note: This section is incomplete for the first review of Draft 1. Many of decisions regarding specific software are detailed in Performance Requirements, Resource Requirements, and Alternative Solutions, above.]

## Overview

### Software design

[Python, SciPy, JQuery, PostgreSQL, PostGIS, Mapnik, TileCache, OpenLayers]

### Data requirements

### Hardware

Should run on any server capable of running Unix/Linux, though Mapnik & TileCache require extra care in terms of CPU power and RAM.

## Testing methods

Django has built-in test suite support[9], based on Python's built-in unit testing support.

[Should also manually test frontend (pages, map rendering) for load times.]

## Scheduling diagram

## Implementation resources & costs

The primary costs underlying NationBrowse relate to the operational costs of the Web server. Additionally, domain name registration fees apply.

The Slicehost server that operates the development prototype costs \$20 per month. The average fee for a .com domain name is about \$10 per year.

# Product implementation

---

In a general sense, our product implementation will consist of a web based application hosted externally on Slicehost — i.e., not on University-operated hardware. This is because the implementation of the GIS software requires a GIS-capable database, which we are not able to install on those systems. The project will be constructed using the Django open source web application framework. Also for this project, Python will be used throughout, especially for settings, files, and data models. We will take advantage of the Javascript library JQuery for the implementation of effects for the front end interface that will make the application look more polished and professional to the end user.

More specifically, the actual user queries from the front end interface will be handled essentially with the Python `eval()` statement. However, extra security measures will have to be considered so the user will not be able to execute arbitrary commands like removing specific files from the server, for example. Luckily, there's a way around this problem. You can pass a second argument to `eval()`, which should contain a dictionary defining the namespace in which the expression is



evaluated. That way, we can define the methods that the user will be able to write in the query string for input. If a method is found in the input sting that is not referenced in the dictionary, it will not be evaluated. There is still the remaining issue of the user being able to put unusually large number arithmetic expressions for the interpreter to evaluate. For example, the expression `eval("() "*8**5)` in Python will cause it to produce a fatal error, causing the interpreters process to be terminated. Currently, we have no solutions to remedy this potential threat to our application.

## Evaluation of product performance

---

The evaluation of our products performance will be through the way of measuring the amount of time it takes for the users query to be fully processed and displayed graphically by the front end of our web application.

## Discussion, future work, and conclusions

---

## References

---

- [1] "Apps For America." 1 April 2009. Retrieved 21 September 2009. <<http://www.sunlightlabs.com/contests/appsforamerica/>>
- [2] "Apps For America 2." 8 August 2009. Retrieved 21 September 2009. <<http://www.sunlightlabs.com/contests/appsforamerica2/>>
- [3] "ThisWeKnow." Retrieved 20 September 2009. <<http://www.thisweknow.org/>>
- [4] "DataMasher." Retrieved 20 September 2009. <<http://www.datamasher.org/>>
- [5] "EveryBlock." Retrieved 20 September 2009. <<http://www.everyblock.com/>>
- [6] Smith, Paul. "Take Control of Your Maps." *A List Apart*. 8 April 2008. Retrieved 20 September 2009. <<http://www.alistapart.com/articles/takecontrolofyourmaps>>
- [7] "Mapping L.A. Neighborhoods." *The Los Angeles Times*. Retrieved 20 September 2009. <<http://projects.latimes.com/mapping-la/neighborhoods/>>
- [8] "Slicehost." Retrieved 23 September 2009. <<http://www.slicehost.com/>>
- [9] "Testing Django applications." Retrieved 22 September 2009. <<http://docs.djangoproject.com/en/dev/topics/testing/>>